

Hands-on Activity 11.2 Classification using Logistic Regression

Name: Xander Sam E. Galapia

Section: CPE22S3

```
pip install ucimlrepo
```

```
Requirement already satisfied: ucimlrepo in /usr/local/lib/python3.10/dist-packages (0.0.6)
```

```
!pip install hvplot
```

```
Requirement already satisfied: hvplot in /usr/local/lib/python3.10/dist-packages (0.9.2)
Requirement already satisfied: bokeh>=1.0.0 in /usr/local/lib/python3.10/dist-packages (from hvplot) (3.3.4)
Requirement already satisfied: colorcet>=2 in /usr/local/lib/python3.10/dist-packages (from hvplot) (3.1.0)
Requirement already satisfied: holoviews>=1.11.0 in /usr/local/lib/python3.10/dist-packages (from hvplot) (1.17.1)
Requirement already satisfied: pandas in /usr/local/lib/python3.10/dist-packages (from hvplot) (2.0.3)
Requirement already satisfied: numpy>=1.15 in /usr/local/lib/python3.10/dist-packages (from hvplot) (1.25.2)
Requirement already satisfied: packaging in /usr/local/lib/python3.10/dist-packages (from hvplot) (24.0)
Requirement already satisfied: panel>=0.11.0 in /usr/local/lib/python3.10/dist-packages (from hvplot) (1.3.8)
Requirement already satisfied: param<3.0,>=1.12.0 in /usr/local/lib/python3.10/dist-packages (from hvplot) (2.1.0)
Requirement already satisfied: Jinja2>=2.9 in /usr/local/lib/python3.10/dist-packages (from bokeh>=1.0.0->hvplot) (3.1.3)
Requirement already satisfied: contourpy>=1 in /usr/local/lib/python3.10/dist-packages (from bokeh>=1.0.0->hvplot) (1.2.1)
Requirement already satisfied: pillow>=7.1.0 in /usr/local/lib/python3.10/dist-packages (from bokeh>=1.0.0->hvplot) (9.4.0)
Requirement already satisfied: PyYAML>=3.10 in /usr/local/lib/python3.10/dist-packages (from bokeh>=1.0.0->hvplot) (6.0.1)
Requirement already satisfied: tornado>=5.1 in /usr/local/lib/python3.10/dist-packages (from bokeh>=1.0.0->hvplot) (6.3.3)
Requirement already satisfied: xyzservices>=2021.09.1 in /usr/local/lib/python3.10/dist-packages (from bokeh>=1.0.0->hvplot) (2024.4.0)
Requirement already satisfied: pyviz-comms>=0.7.4 in /usr/local/lib/python3.10/dist-packages (from holoviews>=1.11.0->hvplot) (3.0.2)
Requirement already satisfied: python-dateutil>=2.8.2 in /usr/local/lib/python3.10/dist-packages (from pandas->hvplot) (2.8.2)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/dist-packages (from pandas->hvplot) (2023.4)
Requirement already satisfied: tzdata>=2022.1 in /usr/local/lib/python3.10/dist-packages (from pandas->hvplot) (2024.1)
Requirement already satisfied: markdown in /usr/local/lib/python3.10/dist-packages (from panel>=0.11.0->hvplot) (3.6)
Requirement already satisfied: markdown-it-py in /usr/local/lib/python3.10/dist-packages (from panel>=0.11.0->hvplot) (3.0.0)
Requirement already satisfied: linkify-it-py in /usr/local/lib/python3.10/dist-packages (from panel>=0.11.0->hvplot) (2.0.3)
Requirement already satisfied: mdit-py-plugins in /usr/local/lib/python3.10/dist-packages (from panel>=0.11.0->hvplot) (0.4.0)
Requirement already satisfied: requests in /usr/local/lib/python3.10/dist-packages (from panel>=0.11.0->hvplot) (2.31.0)
Requirement already satisfied: tqdm>=4.48.0 in /usr/local/lib/python3.10/dist-packages (from panel>=0.11.0->hvplot) (4.66.2)
Requirement already satisfied: bleach in /usr/local/lib/python3.10/dist-packages (from panel>=0.11.0->hvplot) (6.1.0)
Requirement already satisfied: typing-extensions in /usr/local/lib/python3.10/dist-packages (from panel>=0.11.0->hvplot) (4.11.0)
Requirement already satisfied: MarkupSafe>=2.0 in /usr/local/lib/python3.10/dist-packages (from Jinja2>=2.9->bokeh>=1.0.0->hvplot) (2.1.5)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.10/dist-packages (from python-dateutil>=2.8.2->pandas->hvplot) (1.16.0)
Requirement already satisfied: webencodings in /usr/local/lib/python3.10/dist-packages (from bleach->panel>=0.11.0->hvplot) (0.5.1)
Requirement already satisfied: uc-micro-py in /usr/local/lib/python3.10/dist-packages (from linkify-it-py->panel>=0.11.0->hvplot) (1.0.3)
Requirement already satisfied: mdurl~=0.1 in /usr/local/lib/python3.10/dist-packages (from markdown-it-py->panel>=0.11.0->hvplot) (0.1.2)
Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python3.10/dist-packages (from requests->panel>=0.11.0->hvplot) (3.3.2)
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.10/dist-packages (from requests->panel>=0.11.0->hvplot) (3.7)
Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.10/dist-packages (from requests->panel>=0.11.0->hvplot) (2.0.7)
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.10/dist-packages (from requests->panel>=0.11.0->hvplot) (2024.2.2)
```

```
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
%matplotlib inline
```

```
from sklearn.model_selection import train_test_split
from sklearn import metrics
from sklearn.linear_model import LinearRegression
```

```
from ucimlrepo import fetch_ucirepo

# fetch dataset
cervical_cancer_risk_factors = fetch_ucirepo(id=383)


# data (as pandas dataframes)
X = cervical_cancer_risk_factors.data.features
y = cervical_cancer_risk_factors.data.targets

# metadata
print(cervical_cancer_risk_factors.metadata)

# variable information
print(cervical_cancer_risk_factors.variables)
```

17	STDs:syphilis	Feature	Continuous	None
18	STDs:pelvic inflammatory disease	Feature	Continuous	None
19	STDs:genital herpes	Feature	Continuous	None
20	STDs:molluscum contagiosum	Feature	Continuous	None
21	STDs:AIDS	Feature	Continuous	None
22	STDs:HIV	Feature	Continuous	None
23	STDs:Hepatitis B	Feature	Continuous	None
24	STDs:HPV	Feature	Continuous	None
25	STDs: Number of diagnosis	Feature	Integer	None
26	STDs: Time since first diagnosis	Feature	Continuous	None
27	STDs: Time since last diagnosis	Feature	Continuous	None
28	Dx:Cancer	Feature	Integer	None
29	Dx:CIN	Feature	Integer	None
30	Dx:HPV	Feature	Integer	None
31	Dx	Feature	Integer	None
32	Hinselmann	Feature	Integer	None
33	Schiller	Feature	Integer	None
34	Citology	Feature	Integer	None
35	Biopsy	Feature	Integer	None

	description	units	missing_values
0	None	None	no
1	None	None	yes
2	None	None	yes
3	None	None	yes
4	None	None	yes
5	None	None	yes
6	None	None	yes
7	None	None	yes
8	None	None	yes
9	None	None	yes
10	None	None	yes
11	None	None	yes
12	None	None	yes
13	None	None	yes
14	None	None	yes



34NoneNoneno

35NoneNoneno

```
CC = pd.concat([X,y], axis = 1)
CC
```

	Age	Number of sexual partners	First sexual intercourse	Num of pregnancies	Smokes	Smokes (years)	Smokes (packs/year)	Hormonal Contraceptives	Hormonal Contraceptives (years)	IUD	...	STDs: Time since first diagnosis	STDs: Time since last diagnosis	Dx:Cancer	Dx:CIN	Dx:HPV	Dx:Other
0	18	4.0	15.0	1.0	0.0	0.0	0.0	0.0	0.00	0.0	...	NaN	NaN	0	0	0	0
1	15	1.0	14.0	1.0	0.0	0.0	0.0	0.0	0.00	0.0	...	NaN	NaN	0	0	0	0
2	34	1.0	NaN	1.0	0.0	0.0	0.0	0.0	0.00	0.0	...	NaN	NaN	0	0	0	0
3	52	5.0	16.0	4.0	1.0	37.0	37.0	1.0	3.00	0.0	...	NaN	NaN	1	0	1	0
4	46	3.0	21.0	4.0	0.0	0.0	0.0	1.0	15.00	0.0	...	NaN	NaN	0	0	0	0
...
853	34	3.0	18.0	0.0	0.0	0.0	0.0	0.0	0.00	0.0	...	NaN	NaN	0	0	0	0
854	32	2.0	19.0	1.0	0.0	0.0	0.0	1.0	8.00	0.0	...	NaN	NaN	0	0	0	0
855	25	2.0	17.0	0.0	0.0	0.0	0.0	1.0	0.08	0.0	...	NaN	NaN	0	0	0	0
856	33	2.0	24.0	2.0	0.0	0.0	0.0	1.0	0.08	0.0	...	NaN	NaN	0	0	0	0
857	29	2.0	20.0	1.0	0.0	0.0	0.0	1.0	0.50	0.0	...	NaN	NaN	0	0	0	0

```
CC.isna().sum()
```

Age	0
Number of sexual partners	26
First sexual intercourse	7
Num of pregnancies	56
Smokes	13
Smokes (years)	13
Smokes (packs/year)	13
Hormonal Contraceptives	108
Hormonal Contraceptives (years)	108
IUD	117
IUD (years)	117
STDs	105
STDs (number)	105
STDs:condylomatosis	105
STDs:cervical condylomatosis	105
STDs:vaginal condylomatosis	105
STDs:vulvo-perineal condylomatosis	105
STDs:syphilis	105
STDs:pelvic inflammatory disease	105
STDs:genital herpes	105
STDs:molluscum contagiosum	105
STDs:AIDS	105
STDs:HIV	105
STDs:Hepatitis B	105
STDs:HPV	105
STDs: Number of diagnosis	0
STDs: Time since first diagnosis	787
STDs: Time since last diagnosis	787
Dx:Cancer	0
Dx:CIN	0
Dx:HPV	0
Dx	0
Hinselmann	0
Schiller	0
Citology	0
Biopsy	0
dtype: int64	

```
no_val = []
for x in CC.columns:
    if CC[x].isnull().any():
        no_val.append(x)
no_val

['Number of sexual partners',
 'First sexual intercourse',
 'Num of pregnancies',
 'Smokes',
 'Smokes (years)',
 'Smokes (packs/year)',
 'Hormonal Contraceptives',
 'Hormonal Contraceptives (years)',
 'IUD',
 'IUD (years)',
 'STDs',
 'STDs (number)',
 'STDs:condylomatosis',
 'STDs:cervical condylomatosis',
 'STDs:vaginal condylomatosis',
 'STDs:vulvo-perineal condylomatosis',
 'STDs:syphilis',
 'STDs:pelvic inflammatory disease',
 'STDs:genital herpes',
 'STDs:molluscum contagiosum',
 'STDs:AIDS',
 'STDs:HIV',
 'STDs:Hepatitis B',
 'STDs:HPV',
 'STDs: Time since first diagnosis',
 'STDs: Time since last diagnosis']

for x in no_val:
    CC[x].fillna(CC[x].mode()[0], inplace=True)
```

CC.isna().sum()		
Age		0
Number of sexual partners		0
First sexual intercourse		0
Num of pregnancies		0
Smokes		0
Smokes (years)		0
Smokes (packs/year)		0
Hormonal Contraceptives		0
Hormonal Contraceptives (years)		0
IUD		0
IUD (years)		0
STDs		0
STDs (number)		0
STDs:condylomatosis		0
STDs:cervical condylomatosis		0
STDs:vaginal condylomatosis		0
STDs:vulvo-perineal condylomatosis		0
STDs:syphilis		0
STDs:pelvic inflammatory disease		0
STDs:genital herpes		0
STDs:molluscum contagiosum		0
STDs:AIDS		0
STDs:HIV		0
STDs:Hepatitis B		0
STDs:HPV		0
STDs: Number of diagnosis		0
STDs: Time since first diagnosis		0
STDs: Time since last diagnosis		0
Dx:Cancer		0
Dx:CIN		0

Dx:HPV 0
Dx 0
Hinselmann 0
Schiller 0
Citology 0
Biopsy 0
dtype: int64

CC.describe()

	Age	Number of sexual partners	First sexual intercourse	Num of pregnancies	Smokes	Smokes (years)	Smokes (packs/year)	Hormonal Contraceptives	Hormonal Contraceptives (years)	IUD	...	STDs: Time since first diagnosis	STDs: Time since last diagnosis	Dx:
count	858.000000	858.000000	858.000000	858.000000	858.000000	858.000000	858.000000	858.000000	858.000000	858.000000	...	858.000000	858.000000	858.0
mean	26.820513	2.511655	16.979021	2.192308	0.143357	1.201241	0.446278	0.686480	1.972394	0.096737	...	1.425408	1.398601	0.0
std	8.497948	1.644759	2.797653	1.434395	0.350641	4.060623	2.210351	0.464194	3.597888	0.295771	...	2.201550	2.113931	0.0
min	13.000000	1.000000	10.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	...	1.000000	1.000000	0.0
25%	20.000000	2.000000	15.000000	1.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	...	1.000000	1.000000	0.0
50%	25.000000	2.000000	17.000000	2.000000	0.000000	0.000000	0.000000	1.000000	0.250000	0.000000	...	1.000000	1.000000	0.0
75%	32.000000	3.000000	18.000000	3.000000	0.000000	0.000000	0.000000	1.000000	2.000000	0.000000	...	1.000000	1.000000	0.0
max	84.000000	28.000000	32.000000	11.000000	1.000000	37.000000	37.000000	1.000000	30.000000	1.000000	...	22.000000	22.000000	1.0


8 rows × 36 columns

CC.corr()

STDs:condylomatosis	0.016465	0.084794	...	0.384921	0.346300
STDs:cervical condylomatosis	NaN	NaN	NaN	NaN	NaN
STDs:vaginal condylomatosis	0.009505	-0.042120	0.073945	0.002754	0.069651
STDs:vulvo-perineal condylomatosis	-0.011499	0.038992	0.039933	-0.030813	0.058468
STDs:syphilis	0.017457	0.028646	-0.094885	0.150551	0.079358
STDs:pelvic inflammatory disease	0.024854	0.030929	0.000256	-0.052239	-0.013974
STDs:genital herpes	-0.027433	-0.031413	0.024691	-0.028411	-0.013974
STDs:molluscum contagiosum	0.000722	0.030929	-0.011961	0.043074	-0.013974
STDs:AIDS	NaN	NaN	NaN	NaN	NaN
STDs:HIV	0.005009	0.018752	-0.007627	0.014401	0.056151
STDs:Hepatitis B	-0.027433	-0.010633	0.012473	-0.028411	0.083503
STDs:HPV	0.040861	0.014360	0.034938	-0.023343	0.049193
STDs: Number of diagnosis	-0.001606	0.053056	-0.011617	0.039195	0.090725
STDs: Time since first diagnosis	0.089970	0.055506	0.012439	0.073091	0.084156
STDs: Time since last diagnosis	0.106811	0.061756	0.018384	0.084750	0.092835
Dx:Cancer	0.110340	0.023699	0.067996	0.042765	-0.013470
Dx:CIN	0.061443	0.016669	-0.031960	-0.037752	-0.042119
Dx:HPV	0.101722	0.028646	0.044727	0.054111	0.009737
Dx	0.092635	0.024597	0.036664	-0.003034	-0.069396
Hinselmann	-0.003967	-0.039098	-0.015311	0.033987	0.033333
Schiller	0.103283	-0.007230	0.005275	0.077526	0.052028
Citology	-0.016862	0.024067	-0.009594	-0.020131	-0.004639
Biopsy	0.055956	-0.000408	0.008771	0.027959	0.028724

```
CC = CC.drop(['STDs:AIDS',
              'STDs:cervical condylomatosis'], axis = 1)
```

```
plt.figure(figsize =(25,15))
ax = sns.heatmap(CC.corr(), annot = True, cmap = 'Reds')
```




CC.dtypes

Age	int64
Number of sexual partners	float64
First sexual intercourse	float64
Num of pregnancies	float64
Smokes	float64
Smokes (years)	float64
Smokes (packs/year)	float64
Hormonal Contraceptives	float64
Hormonal Contraceptives (years)	float64
IUD	float64
IUD (years)	float64
STDs	float64
STDs (number)	float64
STDs:condylomatosis	float64
STDs:vaginal condylomatosis	float64
STDs:vulvo-perineal condylomatosis	float64
STDs:syphilis	float64
STDs:pelvic inflammatory disease	float64
STDs:genital herpes	float64
STDs:molluscum contagiosum	float64
STDs:HIV	float64
STDs:Hepatitis B	float64
STDs:HPV	float64
STDs: Number of diagnosis	int64
STDs: Time since first diagnosis	float64
STDs: Time since last diagnosis	float64
Dx:Cancer	int64
Dx:CIN	int64
Dx:HPV	int64
Dx	int64
Hinselmann	int64
Schiller	int64
Citology	int64
Biopsy	int64
dtype: object	

CC.head(30)

	Age	sexual partners	sexual intercourse	pregnancies	Smokes	(years)	(packs/year)	Contraceptives	Contraceptives (years)	IUD	...	since first diagnosis	since last diagnosis	Dx:Cancer	Dx:CIN	Dx:HPV
0	18	4.0	15.0	1.0	0.0	0.000000	0.00	0.0	0.00	0.0	...	1.0	1.0	0	0	0
1	15	1.0	14.0	1.0	0.0	0.000000	0.00	0.0	0.00	0.0	...	1.0	1.0	0	0	0
2	34	1.0	15.0	1.0	0.0	0.000000	0.00	0.0	0.00	0.0	...	1.0	1.0	0	0	0
3	52	5.0	16.0	4.0	1.0	37.000000	37.00	1.0	3.00	0.0	...	1.0	1.0	1	0	1
4	46	3.0	21.0	4.0	0.0	0.000000	0.00	1.0	15.00	0.0	...	1.0	1.0	0	0	0
5	42	3.0	23.0	2.0	0.0	0.000000	0.00	0.0	0.00	0.0	...	1.0	1.0	0	0	0
6	51	3.0	17.0	6.0	1.0	34.000000	3.40	0.0	0.00	1.0	...	1.0	1.0	0	0	0
7	26	1.0	26.0	3.0	0.0	0.000000	0.00	1.0	2.00	1.0	...	1.0	1.0	0	0	0
8	45	1.0	20.0	5.0	0.0	0.000000	0.00	0.0	0.00	0.0	...	1.0	1.0	1	0	1
9	44	3.0	15.0	1.0	1.0	1.266973	2.80	0.0	0.00	0.0	...	1.0	1.0	0	0	0
10	44	3.0	26.0	4.0	0.0	0.000000	0.00	1.0	2.00	0.0	...	1.0	1.0	0	0	0
11	27	1.0	17.0	3.0	0.0	0.000000	0.00	1.0	8.00	0.0	...	1.0	1.0	0	0	0
12	45	4.0	14.0	6.0	0.0	0.000000	0.00	1.0	10.00	1.0	...	1.0	1.0	0	0	0
13	44	2.0	25.0	2.0	0.0	0.000000	0.00	1.0	5.00	0.0	...	1.0	1.0	0	0	0
14	43	2.0	18.0	5.0	0.0	0.000000	0.00	0.0	0.00	1.0	...	1.0	1.0	0	0	0
15	40	3.0	18.0	2.0	0.0	0.000000	0.00	1.0	15.00	0.0	...	1.0	1.0	0	0	0
16	41	4.0	21.0	3.0	0.0	0.000000	0.00	1.0	0.25	0.0	...	1.0	1.0	0	0	0
17	43	3.0	15.0	8.0	0.0	0.000000	0.00	1.0	3.00	0.0	...	1.0	1.0	0	0	0
18	42	2.0	20.0	1.0	0.0	0.000000	0.00	1.0	7.00	1.0	...	1.0	1.0	0	0	0
19	40	2.0	27.0	1.0	0.0	0.000000	0.00	0.0	0.00	1.0	...	1.0	1.0	0	0	0
20	43	2.0	18.0	4.0	0.0	0.000000	0.00	1.0	15.00	0.0	...	1.0	1.0	0	0	0
21	41	3.0	17.0	4.0	0.0	0.000000	0.00	1.0	10.00	0.0	...	21.0	21.0	0	0	0
22	40	1.0	18.0	1.0	0.0	0.000000	0.00	1.0	0.25	0.0	...	2.0	2.0	0	0	0
23	40	1.0	20.0	2.0	0.0	0.000000	0.00	1.0	15.00	0.0	...	1.0	1.0	1	0	1
24	40	3.0	15.0	3.0	0.0	0.000000	0.00	1.0	3.00	0.0	...	1.0	1.0	0	0	0
25	44	3.0	19.0	1.0	0.0	0.000000	0.00	0.0	0.00	0.0	...	1.0	1.0	0	0	0
26	39	5.0	23.0	2.0	0.0	0.000000	0.00	0.0	0.00	1.0	...	1.0	1.0	0	0	0
27	39	2.0	17.0	4.0	0.0	0.000000	0.00	0.0	0.00	0.0	...	1.0	1.0	0	0	0
28	37	3.0	24.0	1.0	1.0	3.000000	0.04	0.0	0.00	0.0	...	1.0	1.0	0	0	0
29	37	6.0	26.0	1.0	0.0	0.000000	0.00	1.0	0.25	0.0	...	1.0	1.0	0	0	0

sns.regplot(x = CC['STDs'], y = CC['STDs (number)'])



```
CC.corr()

AX = CC.drop('Dx:Cancer', axis = 1)
AY = CC['Number of sexual partners']

print("X = ",AX.shape, "\nY = ", AY.shape)

X = (858, 33)
Y = (858,)

AX_train, AX_test, AY_train, AY_test = train_test_split(AX, AY, test_size = 0.3, random_state = 101)

AX_train.shape

(600, 33)

AY_train.shape

(600,)

model1 = LinearRegression()

model1.fit(AX_train, AY_train)

LinearRegression()

model1.coef_

array([ 3.33080316e-16,  1.00000000e+00, -2.88423056e-16, -4.00164546e-16,
        -4.64422615e-15,  3.05222863e-16, -2.77110936e-16,  1.09266336e-15,
        -4.14819107e-17, -5.13183030e-15,  5.95855059e-16,  2.63540623e-15,
        -6.75706770e-15,  5.10884151e-15, -5.62762244e-15,  3.30942626e-15,
        -2.20432056e-15,  3.46944695e-18, -1.88096556e-15, -1.73472348e-18,
```

```
1.35785772e-15, -1.22743345e-14, 5.33725057e-15, 4.18966208e-15,
-1.40958362e-16, 2.65004674e-16, 7.93483498e-15, 1.83527187e-15,
-6.24551880e-15, 8.53607961e-16, 9.09769183e-15, 5.06142163e-15,
-1.17494850e-14])
```

```
pd.DataFrame(model1.coef_, AX.columns, columns = ['Coedicients'])
```

number of sexual partners	1.000000e+00
First sexual intercourse	-2.884231e-16
Num of pregnancies	-4.001645e-16
Smokes	-4.644226e-15
Smokes (years)	3.052229e-16
Smokes (packs/year)	-2.771109e-16
Hormonal Contraceptives	1.092663e-15
Hormonal Contraceptives (years)	-4.148191e-17
IUD	-5.131830e-15
IUD (years)	5.958551e-16
STDs	2.635406e-15
STDs (number)	-6.757068e-15
STDs:condylomatosis	5.108842e-15
STDs:vaginal condylomatosis	-5.627622e-15
STDs:vulvo-perineal condylomatosis	3.309426e-15
STDs:syphilis	-2.204321e-15
STDs:pelvic inflammatory disease	3.469447e-18
STDs:genital herpes	-1.880966e-15
STDs:molluscum contagiosum	-1.734723e-18
STDs:HIV	1.357858e-15
STDs:Hepatitis B	-1.227433e-14
STDs:HPV	5.337251e-15
STDs: Number of diagnosis	4.189662e-15
STDs: Time since first diagnosis	-1.409584e-16
STDs: Time since last diagnosis	2.650047e-16
Dx:CIN	7.934835e-15
Dx:HPV	1.835272e-15
Dx	-6.245519e-15
Hinselmann	8.536080e-16
Schiller	9.097692e-15
Citology	5.061422e-15
Biopsy	-1.174949e-14

```
AY_pred = model1.predict(AX_test)
```

```
MAE = metrics.mean_absolute_error(AY_test, AY_pred)
MSE = metrics.mean_squared_error(AY_test, AY_pred)
RMSE = np.sqrt(MSE)
print("MAE: ",MAE)
print("MSE: ",MSE)
print("RMSE: ",RMSE)
```

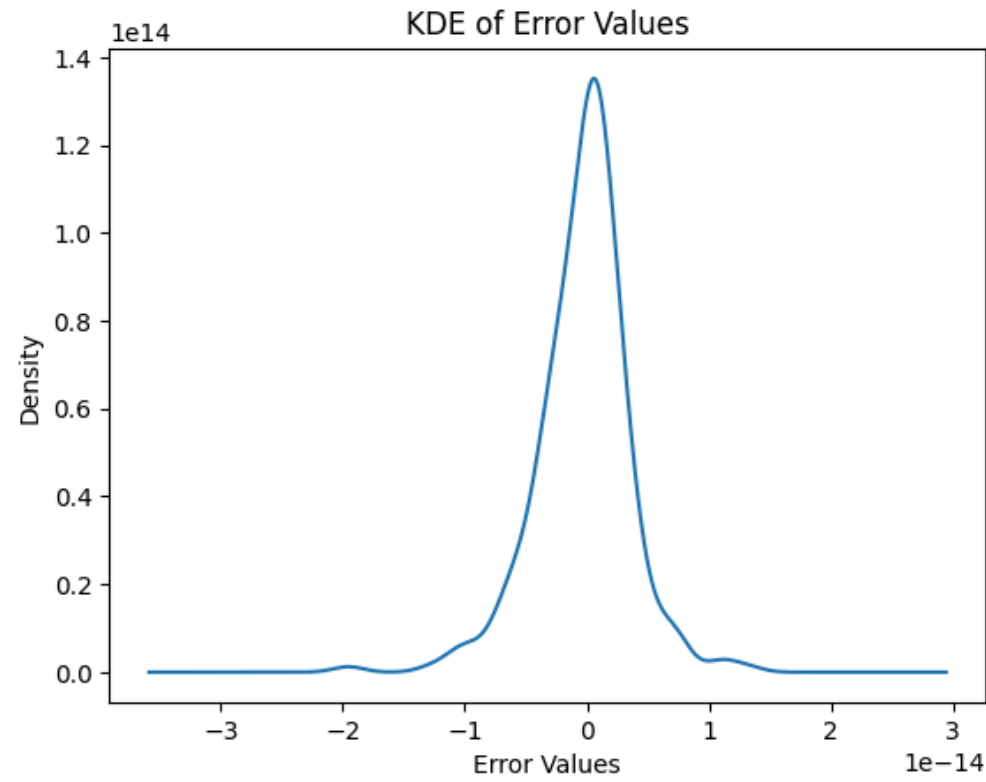
```
MAE:  2.6503347320412168e-15
MSE:  1.4147373762025927e-29
RMSE:  3.7612994778435185e-15
```

```
CC['Dx:Cancer'].mean()
```

```
0.02097902097902098
```

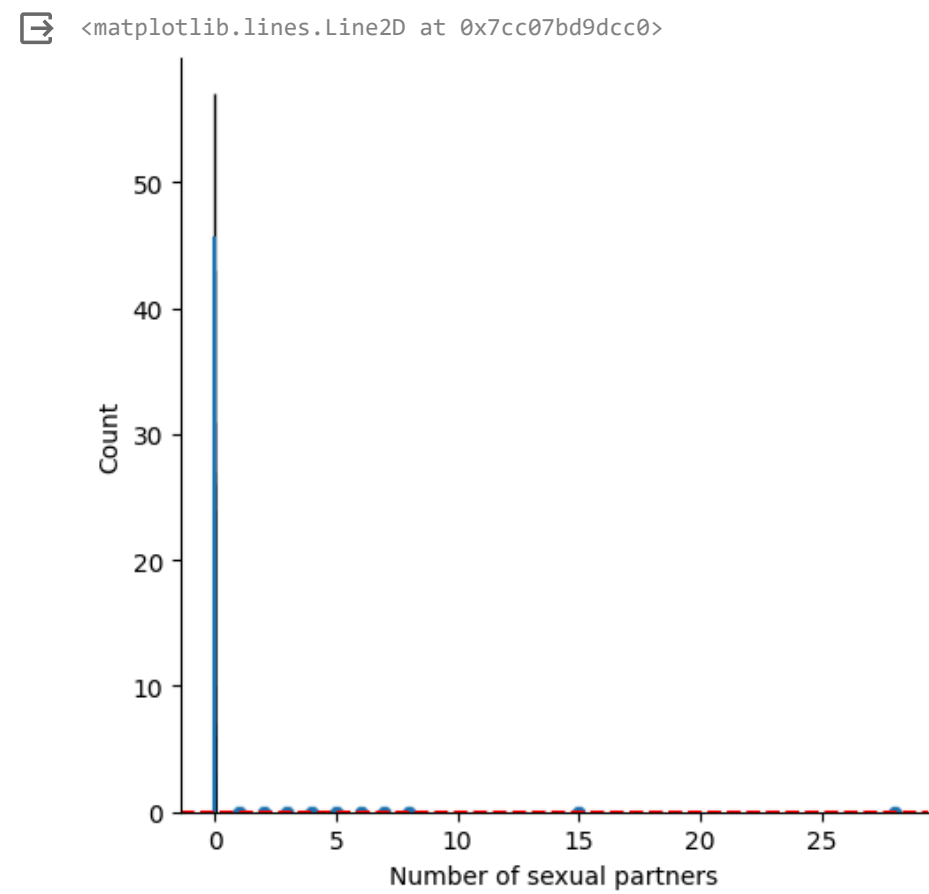
```
test_residual = AY_test - AY_pred
```

```
test_residual.plot(kind='kde')
plt.xlabel('Error Values')
plt.ylabel('Density')
plt.title('KDE of Error Values')
plt.show()
```



```
sns.displot(test_residual, bins = 25, kde = True)
```

```
sns.scatterplot(x = AY_test, y=test_residual)
plt.axhline(y=0, color='red', ls='--')
```



Logistic regression

```
XCan = CC.drop(columns = 'Dx:Cancer')
```

```
YCan = CC['Dx:Cancer']
```

```
XCan_train, XCan_test, YCan_train, YCan_test = train_test_split(XCan, YCan, test_size = 0.3, random_state = 101)
```

```
from sklearn.preprocessing import StandardScaler
```

```
scaler = StandardScaler()
```

```
XCan_train_scaled = scaler.fit_transform(XCan_train)
```

```
XCan_test_scaled = scaler.transform(XCan_test)
```

```
XCan_train_scaled  
  
array([[ 1.14018916,  2.58717906,  3.54934575, ..., -0.30151134,  
        -0.23735633, -0.25635788],  
       [ 0.23572577, -1.08605198,  0.82506721, ..., -0.30151134,  
        -0.23735633, -0.25635788],  
       [-1.23402724, -1.08605198, -1.12084603, ...,  3.31662479,  
        -0.23735633, -0.25635788],  
       ...,  
       [ 0.91407331,  0.38324044,  0.43588457, ..., -0.30151134,  
        -0.23735633, -0.25635788],  
       [ 2.38382633, -0.35140577, -0.34248073, ..., -0.30151134,  
        -0.23735633, -0.25635788],  
       [ 0.57489954,  1.11788665,  0.04670192, ..., -0.30151134,  
        -0.23735633, -0.25635788]])
```

XCan_test_scaled

```
array([[ 0.3487837, -0.3514057,  4.7168937, ..., -0.30151134,
        -0.23735633, -0.25635788],
       [-0.89485347, -0.35140577, -0.73166338, ..., -0.30151134,
        -0.23735633, -0.25635788],
       [-1.23402724, -0.35140577, -0.73166338, ..., -0.30151134,
        4.21307489, -0.25635788],
       ...,
       [ 0.12266785,  0.38324044,  1.60343251, ..., -0.30151134,
        -0.23735633, -0.25635788],
       [-0.66873762, -0.35140577,  0.82506721, ..., -0.30151134,
        -0.23735633, -0.25635788],
       [ 0.00960992,  1.85253285,  0.82506721, ..., -0.30151134,
        4.21307489, -0.25635788]])
```

```
from sklearn.linear_model import LogisticRegression
```

```
from sklearn.metrics import accuracy_score
```

```
log_reg = LogisticRegression(random_state = 0).fit(XCan_train_scaled, YCan_train)
```

```
log_reg.predict(XCan_train_scaled)
```

[illegible]

```
log_reg.score(XCan_train_scaled, YCan_train)
```

1.0

```
log_reg.score(XCan_test_scaled, YCan_test)
```

0.9844961240310077